


# Informed Search




**Georges Sakr**  
ESIB

# Today

- Informed Search
  - Heuristics
  - Greedy Search
  - A\* Search
- Graph Search



## Recap: Search



# Recap: Search

## ■ Search problem:


- States (configurations of the world)
- Actions and costs
- Successor function (world dynamics)
- Start state and goal test

## ■ Search tree:


- Nodes: represent plans for reaching states
- Plans have costs (sum of action costs)

## ■ Search algorithm:

- Systematically builds a search tree
- Chooses an ordering of the fringe (unexplored nodes)
- Optimal: finds least-cost plans





# Uninformed Search



# Uniform Cost Search

- Strategy: expand lowest path cost
- The good: UCS is complete and optimal!
- The bad:
  - Explores options in every “direction”
  - No information about goal location



[Demo: contours UCS empty (L3D1)]

[Demo: contours UCS pacman small maze (L3D3)]

# Informed Search



# Search Heuristics




# Search Heuristics

- A heuristic is:
  - A function that *estimates* how close a state is to a goal
  - Designed for a particular search problem



# Search Heuristics

- A heuristic is:
  - A function that *estimates* how close a state is to a goal
  - Designed for a particular search problem






# Search Heuristics

- A heuristic is:
  - A function that *estimates* how close a state is to a goal
  - Designed for a particular search problem






# Search Heuristics

- A heuristic is:
  - A function that *estimates* how close a state is to a goal
  - Designed for a particular search problem






# Search Heuristics

- A heuristic is:
  - A function that *estimates* how close a state is to a goal
  - Designed for a particular search problem
  - Examples: Manhattan distance, Euclidean distance for pathing






# Search Heuristics

- A heuristic is:
  - A function that *estimates* how close a state is to a goal
  - Designed for a particular search problem
  - Examples: Manhattan distance, Euclidean distance for pathing




# Search Heuristics

- A heuristic is:
  - A function that *estimates* how close a state is to a goal
  - Designed for a particular search problem
  - Examples: Manhattan distance, Euclidean distance for pathing



## Example: Heuristic Function



Straight line distance to ESIB

Bekfaya	366
ESIB	0
Hazmieh	160
Karantina	242
Mansourieh	178
Dekwaneh	77
Jal El Dib	244
Nahr El Mot	241
Broumana	380
Mkales	98
Ain Saade	193
Beit Mery	253
Antelias	329
Baabdat	374

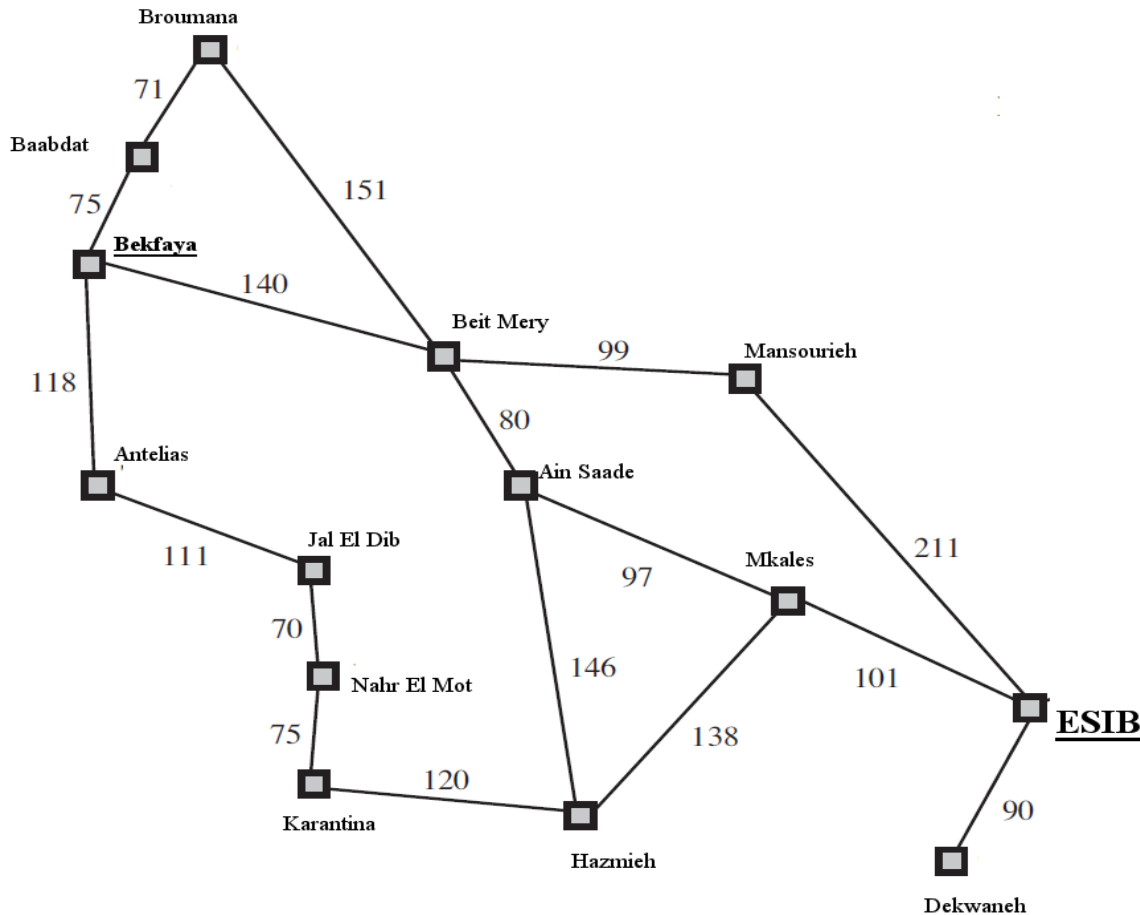
$h(x)$



# Greedy Search



## Example: Heuristic Function




Straight line distance to ESIB

Bekfaya	366
ESIB	0
Hazmieh	160
Karantina	242
Mansourieh	178
Dekwaneh	77
Jal El Dib	244
Nahr El Mot	241
Broumana	380
Mkales	98
Ain Saade	193
Beit Mery	253
Antelias	329
Baabdat	374

$h(x)$


# Greedy Search

Expand the node that seems closest...



# Greedy Search



Expand the node that seems closest...



What can go wrong?

# Greedy Search


Expand the node that seems closest...



What can go wrong?

# Greedy Search

- Strategy: expand a node that you think is closest to a goal state
  - Heuristic: estimate of distance to nearest goal for each state




[Demo: contours greedy empty (L3D1)]

[Demo: contours greedy pacman small maze (L3D4)]

# Greedy Search

- Strategy: expand a node that you think is closest to a goal state
  - Heuristic: estimate of distance to nearest goal for each state
- A common case:
  - Best-first takes you straight to the (wrong) goal




[Demo: contours greedy empty (L3D1)]

[Demo: contours greedy pacman small maze (L3D4)]

# Greedy Search

- Strategy: expand a node that you think is closest to a goal state
  - Heuristic: estimate of distance to nearest goal for each state
- A common case:
  - Best-first takes you straight to the (wrong) goal
- Worst-case: like a badly-guided DFS





[Demo: contours greedy empty (L3D1)]

[Demo: contours greedy pacman small maze (L3D4)]



# Greedy Search

- Strategy: expand a node that you think is closest to a goal state
  - Heuristic: estimate of distance to nearest goal for each state
- A common case:
  - Best-first takes you straight to the (wrong) goal
- Worst-case: like a badly-guided DFS



[Demo: contours greedy empty (L3D1)]

[Demo: contours greedy pacman small maze (L3D4)]

# A\* Search



# A\* Search

# A\* Search



# A\* Search



UCS



# A\* Search



UCS



Greedy

# A\* Search



UCS





Greedy



A\*

## Combining UCS and Greedy





Example: Teg Grenager



## Combining UCS and Greedy



- Uniform-cost orders by path cost, or *backward cost*  $g(n)$



Example: Teg Grenager

## Combining UCS and Greedy

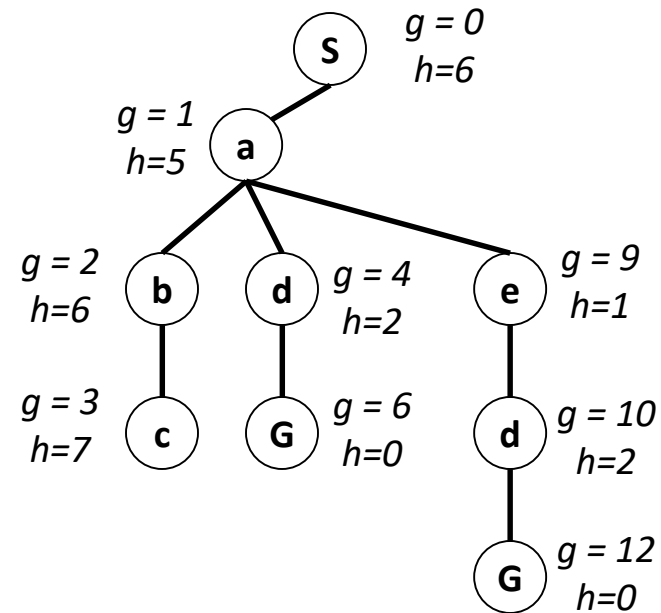

- Uniform-cost orders by path cost, or *backward cost*  $g(n)$



Example: Teg Grenager

## Combining UCS and Greedy



- Uniform-cost orders by path cost, or *backward cost*  $g(n)$



Example: Teg Grenager

## Combining UCS and Greedy

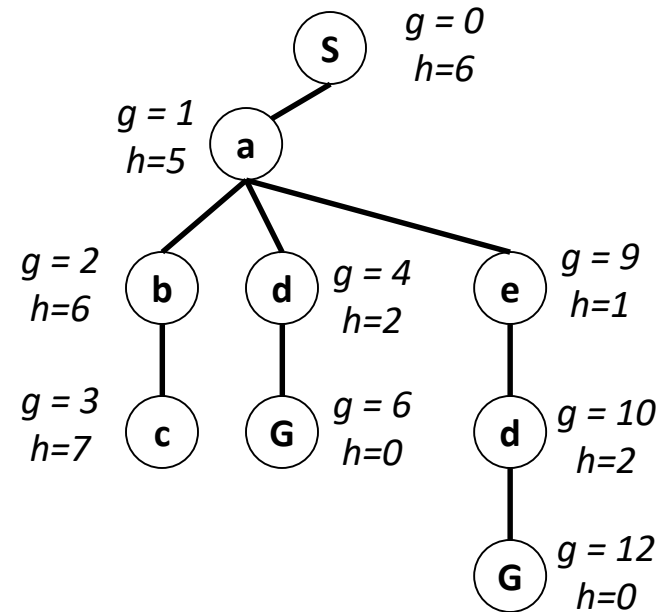

- Uniform-cost orders by path cost, or *backward cost*  $g(n)$
- Greedy orders by goal proximity, or *forward cost*  $h(n)$



Example: Teg Grenager

## Combining UCS and Greedy



- Uniform-cost orders by path cost, or *backward cost*  $g(n)$
- Greedy orders by goal proximity, or *forward cost*  $h(n)$



Example: Teg Grenager

## Combining UCS and Greedy

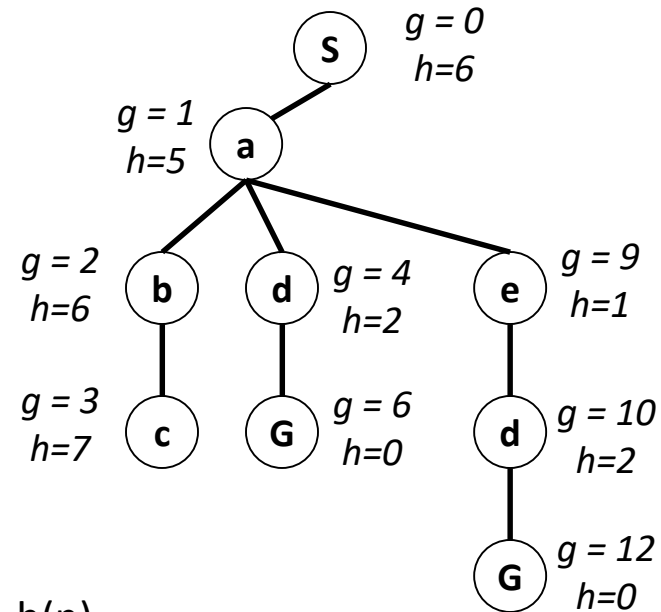

- Uniform-cost orders by path cost, or *backward cost*  $g(n)$
- Greedy orders by goal proximity, or *forward cost*  $h(n)$



Example: Teg Grenager

## Combining UCS and Greedy

- Uniform-cost orders by path cost, or *backward cost*  $g(n)$
- Greedy orders by goal proximity, or *forward cost*  $h(n)$

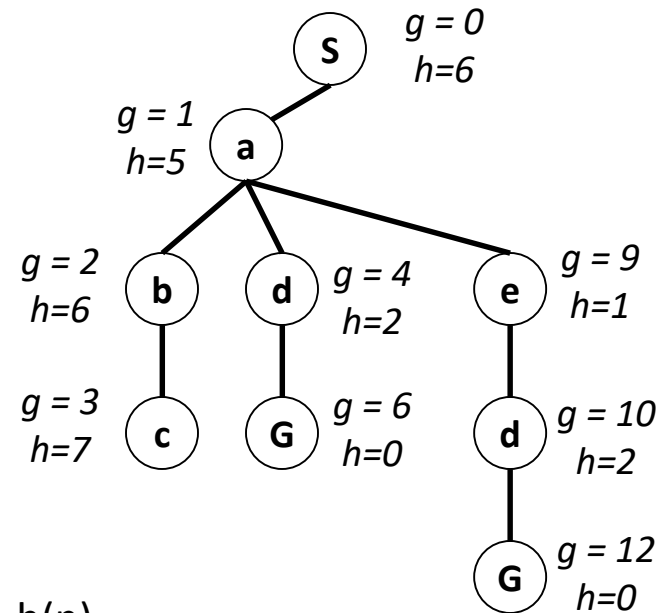



- A\* Search orders by the sum:  $f(n) = g(n) + h(n)$

Example: Teg Grenager

## Combining UCS and Greedy

- Uniform-cost orders by path cost, or *backward cost*  $g(n)$
- Greedy orders by goal proximity, or *forward cost*  $h(n)$




- A\* Search orders by the sum:  $f(n) = g(n) + h(n)$

Example: Teg Grenager




## When should A\* terminate?

- Should we stop when we enqueue a goal?




## When should A\* terminate?

- Should we stop when we enqueue a goal?




- No: only stop when we dequeue a goal

## Is A\* Optimal?




- What went wrong?

## Is A\* Optimal?




- What went wrong?
- Actual bad goal cost < estimated good goal cost

## Is A\* Optimal?




- What went wrong?
- Actual bad goal cost < estimated good goal cost
- We need estimates to be less than actual costs!


# Admissible Heuristics



# Idea: Admissibility



Inadmissible (pessimistic) heuristics break optimality by trapping good plans on the fringe



Admissible (optimistic) heuristics slow down bad plans but never outweigh true costs

# Admissible Heuristics

- A heuristic  $h$  is *admissible* (optimistic) if:

$$0 \leq h(n) \leq h^*(n)$$

where  $h^*(n)$  is the true cost to a nearest goal



# Admissible Heuristics

- A heuristic  $h$  is *admissible* (optimistic) if:

$$0 \leq h(n) \leq h^*(n)$$

where  $h^*(n)$  is the true cost to a nearest goal

- Examples:


# Admissible Heuristics

- A heuristic  $h$  is *admissible* (optimistic) if:

$$0 \leq h(n) \leq h^*(n)$$

where  $h^*(n)$  is the true cost to a nearest goal

- Examples:




# Admissible Heuristics

- A heuristic  $h$  is *admissible* (optimistic) if:

$$0 \leq h(n) \leq h^*(n)$$

where  $h^*(n)$  is the true cost to a nearest goal

- Examples:




# Admissible Heuristics

- A heuristic  $h$  is *admissible* (optimistic) if:

$$0 \leq h(n) \leq h^*(n)$$


where  $h^*(n)$  is the true cost to a nearest goal

- Examples:



- Coming up with admissible heuristics is most of what's involved in using A\* in practice.

# Optimality of A\* Tree Search




## Optimality of A\* Tree Search

Assume:

- A is an optimal goal node
- B is a suboptimal goal node
- $h$  is admissible


Claim:

- A will exit the fringe before B



# Optimality of A\* Tree Search: Blocking


Proof:



## Optimality of A\* Tree Search: Blocking

Proof:

- Imagine B is on the fringe






## Optimality of A\* Tree Search: Blocking

Proof:


- Imagine B is on the fringe



## Optimality of A\* Tree Search: Blocking

Proof:


- Imagine B is on the fringe
- Some ancestor  $n$  of A is on the fringe, too (maybe A!)



## Optimality of A\* Tree Search: Blocking

Proof:


- Imagine B is on the fringe
- Some ancestor  $n$  of A is on the fringe, too (maybe A!)



## Optimality of A\* Tree Search: Blocking

Proof:


- Imagine B is on the fringe
- Some ancestor  $n$  of A is on the fringe, too (maybe A!)
- Claim:  $n$  will be expanded before B



## Optimality of A\* Tree Search: Blocking

Proof:


- Imagine B is on the fringe
- Some ancestor  $n$  of A is on the fringe, too (maybe A!)
- Claim:  $n$  will be expanded before B



## Optimality of A\* Tree Search: Blocking

Proof:

- Imagine B is on the fringe
- Some ancestor  $n$  of A is on the fringe, too (maybe A!)
- Claim:  $n$  will be expanded before B
  1.  $f(n)$  is less or equal to  $f(A)$




## Optimality of A\* Tree Search: Blocking

Proof:

- Imagine  $B$  is on the fringe
- Some ancestor  $n$  of  $A$  is on the fringe, too (maybe  $A$ !)
- Claim:  $n$  will be expanded before  $B$

1.  $f(n)$  is less or equal to  $f(A)$



$$f(n) = g(n) + h(n) \quad \text{Definition of f-cost}$$




## Optimality of A\* Tree Search: Blocking

Proof:

- Imagine B is on the fringe
- Some ancestor  $n$  of A is on the fringe, too (maybe A!)
- Claim:  $n$  will be expanded before B

1.  $f(n)$  is less or equal to  $f(A)$



$$f(n) = g(n) + h(n)$$

Definition of f-cost

$$f(n) \leq g(A)$$


Admissibility of h

## Optimality of A\* Tree Search: Blocking

Proof:

- Imagine B is on the fringe
- Some ancestor  $n$  of A is on the fringe, too (maybe A!)
- Claim:  $n$  will be expanded before B

1.  $f(n)$  is less or equal to  $f(A)$



$$f(n) = g(n) + h(n)$$

Definition of f-cost

$$f(n) \leq g(A)$$

Admissibility of  $h$


$$g(A) = f(A)$$

$h = 0$  at a goal

## Optimality of A\* Tree Search: Blocking

Proof:


- Imagine B is on the fringe
- Some ancestor  $n$  of A is on the fringe, too (maybe A!)
- Claim:  $n$  will be expanded before B
  1.  $f(n)$  is less or equal to  $f(A)$



## Optimality of A\* Tree Search: Blocking

Proof:


- Imagine B is on the fringe
- Some ancestor  $n$  of A is on the fringe, too (maybe A!)
- Claim:  $n$  will be expanded before B
  1.  $f(n)$  is less or equal to  $f(A)$



## Optimality of A\* Tree Search: Blocking

Proof:


- Imagine B is on the fringe
- Some ancestor  $n$  of A is on the fringe, too (maybe A!)
- Claim:  $n$  will be expanded before B
  1.  $f(n)$  is less or equal to  $f(A)$
  2.  $f(A)$  is less than  $f(B)$



## Optimality of A\* Tree Search: Blocking

Proof:


- Imagine B is on the fringe
- Some ancestor  $n$  of A is on the fringe, too (maybe A!)
- Claim:  $n$  will be expanded before B
  1.  $f(n)$  is less or equal to  $f(A)$
  2.  $f(A)$  is less than  $f(B)$



## Optimality of A\* Tree Search: Blocking

Proof:

- Imagine B is on the fringe
- Some ancestor  $n$  of A is on the fringe, too (maybe A!)
- Claim:  $n$  will be expanded before B
  1.  $f(n)$  is less or equal to  $f(A)$
  2.  $f(A)$  is less than  $f(B)$




$$g(A) < g(B)$$

B is suboptimal

## Optimality of A\* Tree Search: Blocking

Proof:

- Imagine B is on the fringe
- Some ancestor  $n$  of A is on the fringe, too (maybe A!)
- Claim:  $n$  will be expanded before B
  1.  $f(n)$  is less or equal to  $f(A)$
  2.  $f(A)$  is less than  $f(B)$



$$g(A) < g(B)$$

$$f(A) < f(B)$$

B is suboptimal


$h = 0$  at a goal



## Optimality of A\* Tree Search: Blocking

Proof:


- Imagine B is on the fringe
- Some ancestor  $n$  of A is on the fringe, too (maybe A!)
- Claim:  $n$  will be expanded before B
  1.  $f(n)$  is less or equal to  $f(A)$
  2.  $f(A)$  is less than  $f(B)$
  3.  $n$  expands before B



## Optimality of A\* Tree Search: Blocking

Proof:


- Imagine B is on the fringe
- Some ancestor  $n$  of A is on the fringe, too (maybe A!)
- Claim:  $n$  will be expanded before B
  1.  $f(n)$  is less or equal to  $f(A)$
  2.  $f(A)$  is less than  $f(B)$
  3.  $n$  expands before B



## Optimality of A\* Tree Search: Blocking

Proof:

- Imagine B is on the fringe
- Some ancestor  $n$  of A is on the fringe, too (maybe A!)
- Claim:  $n$  will be expanded before B
  1.  $f(n)$  is less or equal to  $f(A)$
  2.  $f(A)$  is less than  $f(B)$
  3.  $n$  expands before B




$$f(n) \leq f(A) < f(B)$$

## Optimality of A\* Tree Search: Blocking

Proof:


- Imagine B is on the fringe
- Some ancestor  $n$  of A is on the fringe, too (maybe A!)
- Claim:  $n$  will be expanded before B
  1.  $f(n)$  is less or equal to  $f(A)$
  2.  $f(A)$  is less than  $f(B)$
  3.  $n$  expands before B



## Optimality of A\* Tree Search: Blocking

Proof:


- Imagine B is on the fringe
- Some ancestor  $n$  of A is on the fringe, too (maybe A!)
- Claim:  $n$  will be expanded before B
  1.  $f(n)$  is less or equal to  $f(A)$
  2.  $f(A)$  is less than  $f(B)$
  3.  $n$  expands before B
- All ancestors of A expand before B



## Optimality of A\* Tree Search: Blocking

Proof:


- Imagine B is on the fringe
- Some ancestor  $n$  of A is on the fringe, too (maybe A!)
- Claim:  $n$  will be expanded before B
  1.  $f(n)$  is less or equal to  $f(A)$
  2.  $f(A)$  is less than  $f(B)$
  3.  $n$  expands before B
- All ancestors of A expand before B
- A expands before B



## Optimality of A\* Tree Search: Blocking

Proof:

- Imagine B is on the fringe
- Some ancestor  $n$  of A is on the fringe, too (maybe A!)
- Claim:  $n$  will be expanded before B
  1.  $f(n)$  is less or equal to  $f(A)$
  2.  $f(A)$  is less than  $f(B)$
  3.  $n$  expands before B
- All ancestors of A expand before B
- A expands before B
- A\* search is optimal




# Properties of $A^*$




# Properties of A\*

Uniform-Cost





A\*



## UCS vs A\* Contours

- Uniform-cost expands equally in all “directions”
- A\* expands mainly toward the goal, but does hedge its bets to ensure optimality



[Demo: contours UCS / greedy / A\* empty (L3D1)]

[Demo: contours A\* pacman small maze (L3D5)]

# Comparison



Greedy



Uniform Cost



A\*

# A\* Applications



# A\* Applications


- Video games
- Pathing / routing problems
- Resource planning problems
- Robot motion planning
- Language analysis
- Machine translation
- Speech recognition
- ...



[Demo: UCS / A\* pacman tiny maze (L3D6,L3D7)]

[Demo: guess algorithm Empty Shallow/Deep (L3D8)]

# Creating Heuristics




# Creating Admissible Heuristics

- Most of the work in solving hard search problems optimally is in coming up with admissible heuristics

# Creating Admissible Heuristics



- Most of the work in solving hard search problems optimally is in coming up with admissible heuristics
- Often, admissible heuristics are solutions to *relaxed problems*, where new actions are available






# Creating Admissible Heuristics

- Most of the work in solving hard search problems optimally is in coming up with admissible heuristics
- Often, admissible heuristics are solutions to *relaxed problems*, where new actions are available




# Creating Admissible Heuristics

- Most of the work in solving hard search problems optimally is in coming up with admissible heuristics
- Often, admissible heuristics are solutions to *relaxed problems*, where new actions are available




- Inadmissible heuristics are often useful too



## Example: 8 Puzzle

7	2	4
5		6
8	3	1

Start State



Actions


	1	2
3	4	5
6	7	8

Goal State


- What are the states?
- How many states?
- What are the actions?
- How many successors from the start state?
- What should the costs be?

## 8 Puzzle I

- Heuristic: Number of tiles misplaced




Start State




Goal State

## 8 Puzzle I

- Heuristic: Number of tiles misplaced
- Why is it admissible?




Start State




Goal State

## 8 Puzzle I

- Heuristic: Number of tiles misplaced
- Why is it admissible?
- $h(\text{start}) =$




Start State




Goal State

## 8 Puzzle I

- Heuristic: Number of tiles misplaced
- Why is it admissible?
- $h(\text{start}) =$



Start State




Goal State

Average nodes expanded when the optimal path has...			
	...4 steps	...8 steps	...12 steps
UCS	112	6,300	$3.6 \times 10^6$
TILES	13	39	227


Statistics from Andrew Moore

## 8 Puzzle I

- Heuristic: Number of tiles misplaced
- Why is it admissible?
- $h(\text{start}) =$
- This is a *relaxed-problem* heuristic



Start State



Goal State


Average nodes expanded when the optimal path has...			
	...4 steps	...8 steps	...12 steps
UCS	112	6,300	$3.6 \times 10^6$
TILES	13	39	227

Statistics from Andrew Moore




## 8 Puzzle I


- Heuristic: Number of tiles misplaced
- Why is it admissible?
- $h(\text{start}) =$
- This is a *relaxed-problem* heuristic



Start State



Goal State




Average nodes expanded when the optimal path has...			
	...4 steps	...8 steps	...12 steps
UCS	112	6,300	$3.6 \times 10^6$
TILES	13	39	227


Statistics from Andrew Moore

## 8 Puzzle II

- What if we had an easier 8-puzzle where any tile could slide any direction at any time, ignoring other tiles?




Start State




Goal State

## 8 Puzzle II

- What if we had an easier 8-puzzle where any tile could slide any direction at any time, ignoring other tiles?
- Total *Manhattan* distance




Start State




Goal State

## 8 Puzzle II

- What if we had an easier 8-puzzle where any tile could slide any direction at any time, ignoring other tiles?
- Total *Manhattan* distance
- Why is it admissible?




Start State




Goal State

## 8 Puzzle II

- What if we had an easier 8-puzzle where any tile could slide any direction at any time, ignoring other tiles?
- Total *Manhattan* distance
- Why is it admissible?
- $h(\text{start}) =$




Start State




Goal State

## 8 Puzzle II

- What if we had an easier 8-puzzle where any tile could slide any direction at any time, ignoring other tiles?
- Total *Manhattan* distance
- Why is it admissible?
- $h(\text{start}) = 3 + 1 + 2 + \dots = 18$




Start State




Goal State

## 8 Puzzle II

- What if we had an easier 8-puzzle where any tile could slide any direction at any time, ignoring other tiles?



Start State



Goal State

- Total *Manhattan* distance
- Why is it admissible?
- $h(\text{start}) = 3 + 1 + 2 + \dots = 18$

Average nodes expanded when the optimal path has...			
	...4 steps	...8 steps	...12 steps
TILES	13	39	227
MANHATTAN	12	25	73

## 8 Puzzle III

- How about using the *actual cost* as a heuristic?
  - Would it be admissible?
  - Would we save on nodes expanded?
  - What's wrong with it?





## 8 Puzzle III

- How about using the *actual cost* as a heuristic?
  - Would it be admissible?
  - Would we save on nodes expanded?
  - What's wrong with it?
- With A\*: a trade-off between quality of estimate and work per node

## 8 Puzzle III

- How about using the *actual cost* as a heuristic?

- Would it be admissible?
- Would we save on nodes expanded?
- What's wrong with it?





- With A\*: a trade-off between quality of estimate and work per node

## 8 Puzzle III

- How about using the *actual cost* as a heuristic?

- Would it be admissible?
- Would we save on nodes expanded?
- What's wrong with it?




- With A\*: a trade-off between quality of estimate and work per node

- As heuristics get closer to the true cost, you will expand fewer nodes but usually do more work per node to compute the heuristic itself


# Semi-Lattice of Heuristics

# Trivial Heuristics, Dominance

- Dominance:  $h_a \geq h_c$  if  
 $\forall n : h_a(n) \geq h_c(n)$
- Heuristics form a semi-lattice:
  - Max of admissible heuristics is admissible  
 $h(n) = \max(h_a(n), h_b(n))$
- Trivial heuristics
  - Bottom of lattice is the zero heuristic (what does this give us?)
  - Top of lattice is the exact heuristic





# Graph Search





## Tree Search: Extra Work!

- Failure to detect repeated states can cause exponentially more work.



## Tree Search: Extra Work!


- Failure to detect repeated states can cause exponentially more work.






## Graph Search

- In BFS, for example, we shouldn't bother expanding the circled nodes (why?)



## Graph Search

- In BFS, for example, we shouldn't bother expanding the circled nodes (why?)



# Graph Search

- Idea: never **expand** a state twice

# Graph Search

- Idea: never **expand** a state twice
- How to implement:
  - Tree search + set of expanded states ("closed set")
  - Expand the search tree node-by-node, but...
  - Before expanding a node, check to make sure its state has never been expanded before
  - If not new, skip it, if new add to closed set

# Graph Search

- Idea: never **expand** a state twice
- How to implement:
  - Tree search + set of expanded states ("closed set")
  - Expand the search tree node-by-node, but...
  - Before expanding a node, check to make sure its state has never been expanded before
  - If not new, skip it, if new add to closed set
- Important: **store the closed set as a set**, not a list

# Graph Search


- Idea: never **expand** a state twice
- How to implement:
  - Tree search + set of expanded states ("closed set")
  - Expand the search tree node-by-node, but...
  - Before expanding a node, check to make sure its state has never been expanded before
  - If not new, skip it, if new add to closed set
- Important: **store the closed set as a set**, not a list
- Can graph search wreck completeness? Why/why not?

# Graph Search

- Idea: never **expand** a state twice
- How to implement:
  - Tree search + set of expanded states ("closed set")
  - Expand the search tree node-by-node, but...
  - Before expanding a node, check to make sure its state has never been expanded before
  - If not new, skip it, if new add to closed set
- Important: **store the closed set as a set**, not a list
- Can graph search wreck completeness? Why/why not?
- How about optimality?

## A\* Graph Search Gone Wrong?


State space graph






## A\* Graph Search Gone Wrong?

State space graph



## A\* Graph Search Gone Wrong?

State space graph




Search tree


S (0+2)

## A\* Graph Search Gone Wrong?

State space graph




Search tree




## A\* Graph Search Gone Wrong?

State space graph




Search tree




## A\* Graph Search Gone Wrong?

State space graph




Search tree




## A\* Graph Search Gone Wrong?

State space graph




Search tree




## A\* Graph Search Gone Wrong?

State space graph



Search tree



# Consistency of Heuristics




## Consistency of Heuristics

- Main idea: estimated heuristic costs  $\leq$  actual costs

## Consistency of Heuristics


- Main idea: estimated heuristic costs  $\leq$  actual costs
  - Admissibility: heuristic cost  $\leq$  actual cost to goal

## Consistency of Heuristics



- Main idea: estimated heuristic costs  $\leq$  actual costs
  - Admissibility: heuristic cost  $\leq$  actual cost to goal

## Consistency of Heuristics




- Main idea: estimated heuristic costs  $\leq$  actual costs

- Admissibility: heuristic cost  $\leq$  actual cost to goal

$$h(A) \leq \text{actual cost from A to G}$$

## Consistency of Heuristics




$h=4$

- Main idea: estimated heuristic costs  $\leq$  actual costs


- Admissibility: heuristic cost  $\leq$  actual cost to goal

$$h(A) \leq \text{actual cost from A to G}$$

## Consistency of Heuristics



$h=4$



- Main idea: estimated heuristic costs  $\leq$  actual costs

- Admissibility: heuristic cost  $\leq$  actual cost to goal


$$h(A) \leq \text{actual cost from A to G}$$

## Consistency of Heuristics

- Main idea: estimated heuristic costs  $\leq$  actual costs

- Admissibility: heuristic cost  $\leq$  actual cost to goal

$$h(A) \leq \text{actual cost from A to G}$$




## Consistency of Heuristics

- Main idea: estimated heuristic costs  $\leq$  actual costs


- Admissibility: heuristic cost  $\leq$  actual cost to goal

$$h(A) \leq \text{actual cost from A to G}$$






## Consistency of Heuristics




- Main idea: estimated heuristic costs  $\leq$  actual costs
  - Admissibility: heuristic cost  $\leq$  actual cost to goal  
 $h(A) \leq \text{actual cost from A to G}$
  - Consistency: heuristic “arc” cost  $\leq$  actual cost for each arc

## Consistency of Heuristics




- Main idea: estimated heuristic costs  $\leq$  actual costs
- Admissibility: heuristic cost  $\leq$  actual cost to goal  
 $h(A) \leq \text{actual cost from A to G}$
- Consistency: heuristic “arc” cost  $\leq$  actual cost for each arc

## Consistency of Heuristics




- Main idea: estimated heuristic costs  $\leq$  actual costs
  - Admissibility: heuristic cost  $\leq$  actual cost to goal  
 $h(A) \leq \text{actual cost from A to G}$
  - Consistency: heuristic “arc” cost  $\leq$  actual cost for each arc

## Consistency of Heuristics



- Main idea: estimated heuristic costs  $\leq$  actual costs
  - Admissibility: heuristic cost  $\leq$  actual cost to goal  
 $h(A) \leq \text{actual cost from A to G}$
  - Consistency: heuristic “arc” cost  $\leq$  actual cost for each arc

## Consistency of Heuristics




- Main idea: estimated heuristic costs  $\leq$  actual costs
  - Admissibility: heuristic cost  $\leq$  actual cost to goal  
 $h(A) \leq \text{actual cost from A to G}$
  - Consistency: heuristic “arc” cost  $\leq$  actual cost for each arc  
 $h(A) - h(C) \leq \text{cost}(A \text{ to } C)$

## Consistency of Heuristics


- Main idea: estimated heuristic costs  $\leq$  actual costs

- Admissibility: heuristic cost  $\leq$  actual cost to goal  
 $h(A) \leq \text{actual cost from A to G}$
- Consistency: heuristic “arc” cost  $\leq$  actual cost for each arc




$$h(A) - h(C) \leq \text{cost}(A \text{ to } C)$$

## Consistency of Heuristics



- Main idea: estimated heuristic costs  $\leq$  actual costs
- Admissibility: heuristic cost  $\leq$  actual cost to goal  
 $h(A) \leq \text{actual cost from A to G}$
- Consistency: heuristic “arc” cost  $\leq$  actual cost for each arc  
 $h(A) - h(C) \leq \text{cost}(A \text{ to } C)$
- Consequences of consistency:


## Consistency of Heuristics



- Main idea: estimated heuristic costs  $\leq$  actual costs
  - Admissibility: heuristic cost  $\leq$  actual cost to goal  
 $h(A) \leq \text{actual cost from A to G}$
  - Consistency: heuristic “arc” cost  $\leq$  actual cost for each arc  
 $h(A) - h(C) \leq \text{cost}(A \text{ to } C)$
- Consequences of consistency:
  - The f value along a path never decreases



## Consistency of Heuristics




- Main idea: estimated heuristic costs  $\leq$  actual costs
  - Admissibility: heuristic cost  $\leq$  actual cost to goal  

$$h(A) \leq \text{actual cost from A to G}$$
  - Consistency: heuristic “arc” cost  $\leq$  actual cost for each arc  

$$h(A) - h(C) \leq \text{cost}(A \text{ to } C)$$
- Consequences of consistency:
  - The f value along a path never decreases  


$$h(A) \leq \text{cost}(A \text{ to } C) + h(C)$$

## Consistency of Heuristics




- Main idea: estimated heuristic costs  $\leq$  actual costs
  - Admissibility: heuristic cost  $\leq$  actual cost to goal
 
$$h(A) \leq \text{actual cost from A to G}$$
  - Consistency: heuristic “arc” cost  $\leq$  actual cost for each arc
 
$$h(A) - h(C) \leq \text{cost}(A \text{ to } C)$$
- Consequences of consistency:
  - The f value along a path never decreases
 
$$h(A) \leq \text{cost}(A \text{ to } C) + h(C)$$
  - A\* graph search is optimal

# Optimality of A\* Graph Search




## Optimality of A\* Graph Search

- Sketch: consider what A\* does with a consistent heuristic:
  - Fact 1: In tree search, A\* expands nodes in increasing total f value (f-contours)
  - Fact 2: For every state s, nodes that reach s optimally are expanded before nodes that reach s suboptimally
  - Result: A\* graph search is optimal



# Optimality

- **Tree search:**
  - A\* is optimal if heuristic is admissible
  - UCS is a special case ( $h = 0$ )
- **Graph search:**
  - A\* optimal if heuristic is consistent
  - UCS optimal ( $h = 0$  is consistent)
- **Consistency implies admissibility**
- **In general, most natural admissible heuristics tend to be consistent, especially if from relaxed problems**



## A\*: Summary



## A\*: Summary

- A\* uses both backward costs and (estimates of) forward costs
- A\* is optimal with admissible / consistent heuristics
- Heuristic design is key: often use relaxed problems

